

Property-Based Testing for the Robot Operating System

André Santos
INESC TEC & University of Minho
Braga, Portugal
andre.f.santos@inesctec.pt

Alcino Cunha
INESC TEC & University of Minho
Braga, Portugal
alcino.cunha@inesctec.pt

Nuno Macedo
INESC TEC & University of Minho
Braga, Portugal
nuno.m.macedo@inesctec.pt

ABSTRACT

The Robot Operating System (ROS) is an open source framework for the development of robotic software, in which a typical system consists of multiple processes communicating under a publisher-subscriber architecture. A great deal of development time goes into orchestration and making sure that the communication interfaces comply with the expected contracts (e.g. receiving a message leads to the publication of another message). Orchestration mistakes are only detected during runtime, stressing the importance of component and integration testing in the verification process. Property-based Testing is fitting in this context, since it is based on the specification of contracts and treats tested components as black boxes, but there is no support for it in ROS. In this paper, we present a first approach towards automatic generation of test scripts for property-based testing of various configurations of a ROS system.

CCS CONCEPTS

• **Software and its engineering** → **Software testing and debugging**; *Publish-subscribe / event-based architectures*; • **Computer systems organization** → *Robotics*;

KEYWORDS

Software Testing, Test Automation, Property-based Testing, Robot Operating System

ACM Reference Format:

André Santos, Alcino Cunha, and Nuno Macedo. 2018. Property-Based Testing for the Robot Operating System. In *Proceedings of the 9th ACM SIGSOFT International Workshop on Automating TEST Case Design, Selection, and Evaluation (A-TEST '18)*, November 5, 2018, Lake Buena Vista, FL, USA. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/3278186.3278195>

1 INTRODUCTION

Software testing has become a staple in most quality assurance processes, providing valuable feedback regarding the correctness of the software under test. In particular, testing is often a mandatory step in the verification of safety-critical systems [10].

With the recent developments in robotics, robots are taking their place in various safety-critical application domains, such as health,

industry and agriculture. As their capabilities increase, so does the complexity of the software behind these systems. Considering the possible dangers of human-robot interaction, and how expensive the systems themselves are, ensuring that they are correct and well tested is a necessity. However, developing a fairly complete test suite for a robot, covering a wide variety of scenarios, can be challenging, especially at the level of integration and system testing, due to the heavy coupling of these systems on input and output (sensors and actuators).

The Robot Operating System (ROS)¹ [11] is one of the most popular open source frameworks for the development of robotics systems, with thousands of users worldwide. It provides middleware, libraries and tools that aim to shorten development time and encourage re-use of existing components. A typical ROS system is a set of independent processes (called *nodes*) communicating with each other through message-passing. Most of the time, nodes communicate asynchronously in a publisher-subscriber fashion (called *topics*), although a client-server model is available (called *services*).

Due to the distributed and re-usable nature of ROS components, part of a ROS developer's job is to integrate nodes, using configuration files and name aliases, to ensure that topics and services match. Component integration is a challenging task [4] that is very prone to human error, and a common source of bugs. Such integration bugs tend to be detected early on, as developers notice when a node is not receiving expected messages. Nonetheless, these issues are only detected manually with the assistance of runtime inspection tools. In this regard, automated testing would decrease development time, besides providing a systematic way to make sure that the whole configuration adheres to an intended architecture.

There is some support in ROS for automated testing² using popular testing libraries, such as Google's *gtest* for C++ code and *unittest* for Python code, but this is mostly appropriate for library unit tests and node unit tests (testing the ROS interface of a single node). ROS also provides simulation environments and message replay tools, which alleviate hardware dependencies when testing.

Many desirable safety properties in a ROS system apply at the interface and integration levels, and, thus, unit testing does not suffice. For instance, consider a simple system in which a node handles safety behaviours, while another node publishes sensor readings and subscribes to actuator commands. This is the case, for instance, with Kobuki³, a popular ROS robot used in research and education. A desirable property for this system would be to ensure that the safety controller node publishes a command (e.g. stopping) after receiving a sensor message signalling a bump into an obstacle, and that this message is received by the base node, to relay it to the actuators. These properties can be further refined by adding

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
A-TEST '18, November 5, 2018, Lake Buena Vista, FL, USA

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-6053-1/18/11...\$15.00

<https://doi.org/10.1145/3278186.3278195>

¹<http://www.ros.org/>

²<http://wiki.ros.org/Quality/Tutorials/UnitTesting>

³<http://kobuki.yujinrobot.com/>

timing constraints (e.g. the stop command being published within, at most, one second after the bump).

Such properties can be specified and tested with manually crafted test scripts, but this can prove to be cumbersome with the default testing framework. A more efficient strategy is to exploit Property-based Testing (PBT), a property-oriented testing method in which common approaches use seemingly random input generators in a systematic way, attempting to falsify properties specified by the developers. Using random inputs may lead to redundant tests, but, on the other hand, they have the potential to uncover defects that a test developer would not come up with.

Our approach adapts this idea to the context of a ROS system. Simply put, we take a set of nodes, ranging from a single node to an entire application, and we consider it as a black box. The inputs for this black box are open subscribed topics (topics where there are no active publishers) and the outputs are open published topics (topics where there are no active subscribers). Then we use Hypothesis⁴ [8], a PBT library, to find a sequence of ROS messages that either crashes the configuration under test, or falsifies a specified property. Due to the random nature of PBT, this approach is most suitable for pure software nodes or hardware controllers using simulation.

Still, writing such tests for various ROS configurations would be as error prone as orchestrating the systems themselves, since node and topic names must be specified. Ideally, this task should be automated, so that, for instance, typing mistakes are completely avoided. As such, inspired by Model-based Testing approaches, we propose an automatic test generation method from configuration models extracted using HAROS⁵ [13], a static analysis framework for ROS applications. Our contribution, thus, is a prototype that is capable of taking models of ROS configurations and generating customisable, property-based test scripts for said configurations.

Throughout the paper, we will use a fictitious mobile ROS robot, called FictiBot⁶, as our running example. This robot is essentially composed of two ROS nodes: `fictibase`, a low-level driver node that directly controls the robot's sensors and actuators; and `ficticontrol`, a higher-level controller node that generates a random robot command based on sensor readings. When both nodes are properly integrated, the driver feeds the controller by publishing sensor information, while the controller feeds the driver back by publishing robot commands.

In the remainder of this paper, we start, in Section 2, by providing necessary background for understanding the problem and our contribution. Then we present our prototype test generator, in Section 3, going into detail on the test generation process, from HAROS models to a concrete test script, but also describing the inner workings of the test script itself. We compare our approach to other relevant related work in Section 4. Finally, we wrap up with a few conclusions and our directions of future work, in Section 5.

2 BACKGROUND

2.1 Property-based Testing and Hypothesis

Property-based Testing is a testing technique in which the testing process is driven by the specification of generic properties that a

system should satisfy [7]. Properties are often assertions on how outputs relate to inputs, and often fall into common patterns, such as how two procedures achieve an equivalent result. A typical example is that applying reverse on the reverse of a list returns the original list, regardless of the list. This immediately contrasts with typical example-based testing methods, where the goal is to verify that the system behaves well for a number of crafted scenarios. An advantage of PBT is that it requires a more formal reasoning about the tested systems than traditional testing methods do. The formulation of specifications, in turn, helps in documenting and building hierarchical models of the system behaviour [2].

In practice, PBT libraries take the property-based test cases and execute them repeatedly, in an attempt to find a counterexample for the specification. Inputs are automatically selected from a large corpus, possibly dynamically generated, and possibly selected at random. Going back to the reverse example, that property would be tested with lists of varying length and contents, ranging from an empty list up to lists of a certain length limit.

PBT is most often applied in unit testing, although it has been applied in testing of web services, for instance. It became popular as a testing method with QuickCheck⁷ [5] for the Haskell programming language. It fits very naturally with the pure functional programming style of Haskell, although modern PBT libraries [2, 5, 8] (QuickCheck included) are able to handle stateful systems and software with side effects just as well. Hypothesis is an example of such a PBT library, mostly for the Python programming language.

Stateful systems can be tested with Hypothesis using a construct called a `RuleBasedStateMachine`. Test developers extend the `RuleBasedStateMachine` class to add their own internal state, to define set up and tearing down of class instances, and to define allowed operations (state transitions, called *rules* in Hypothesis). Additionally, rules can be decorated with preconditions, and invariants can also be specified. Properties in this case are regular Python assertions, written within rules or invariant definitions. The test execution consists of Hypothesis repeatedly creating instances of the state machine, and executing a sequence of rules with generated inputs.

An example of stateful testing in Hypothesis is provided in Figure 1. This is a naïve example, asserting that trying to remove a random integer from a list of integers will result in a list of the same length or shorter. This property would be correct, in theory, but calling `remove` in Python with an element that is not present in the list results in an error. To fix this example, either the call to `remove` must be guarded with a conditional, or the error handled.

2.2 The Robot Operating System

ROS is a multi-language framework, where nodes written in different languages – the two most common ones being C++ and Python – can interact seamlessly. A typical ROS system consists of a peer-to-peer network of nodes, exchanging messages using *topics* for a publisher-subscriber model, or *services* for a client-server model. All nodes using a topic or service should exchange messages of the same type. ROS provides a few basic message types, but users can define their own using a message definition language. During the

⁴<https://github.com/HypothesisWorks/hypothesis>

⁵<https://github.com/git-afsantos/haros>

⁶Source code available at https://github.com/git-afsantos/haros_tutorials

⁷<http://hackage.haskell.org/package/QuickCheck>

```

1 class ListMachine(RuleBasedStateMachine):
2     def __init__(self):
3         super(ListMachine, self).__init__()
4         self.state = []
5     @rule(value=integers())
6     def append(self, value):
7         self.state.append(value)
8     @rule(value=integers())
9     def remove(self, value):
10        previous_length = len(self.state)
11        self.state.remove(value)
12        assert len(self.state) <= previous_length

```

Figure 1: Minimal Hypothesis example for stateful testing.




Figure 2: ROS Computation Graph of FictiBot. Circles represent nodes, rectangles represent topics. Outgoing arrows represent publications, incoming arrows represent subscriptions.

project's build phase, ROS tools generate source code for message types in all target languages.

There is a special node in all ROS systems, called the *ROS Master*, which is always started by default. Its purpose is to help set up the network, by providing nodes with peer discovery, and to hold a shared key-value store where nodes can read and write key-value pairs (called *parameters*) at runtime. Collectively, nodes, topics, services and parameters are called *resources*, and the network of ROS resources is called the *ROS Computation Graph*⁸. Figure 2 shows a diagram of the Computation Graph for our FictiBot running example. Note that the ROS Master is omitted from the diagram.

Computation Graphs are dynamic, meaning that resources can join and leave at any time. Developers use XML-like files, called *launch files*, to deploy sets of nodes and parameters into a Computation Graph (or to create a new one), using a tool called *roslaunch*. Launch files also provide mechanisms to access environment information, group resources under a namespace and deploy resources conditionally. Another core feature of launch files is a name aliasing mechanism, called *remappings* or *remaps*, which is fundamental to orchestrate and re-use nodes. For instance, if a node is configured with a remapping $A \rightarrow B$, it would be transparently forwarded to a resource named B upon requiring a resource named A . Verification of launch files is mostly limited to syntax and lookup of node executables, which makes component integration all the more challenging. Figure 3 shows an example launch file, used to deploy the driver and controller nodes of our running example.

⁸<http://wiki.ros.org/ROS/Concepts>

```

1 <launch>      <!-- File "minimal.launch" -->
2 <node name="fictibase" pkg="fictibot_drivers"
3   type="fictibot_driver" />
4 <node name="ficticontrol" pkg="fictibot_controller"
5   type="fictibot_controller" />
6 </launch>

```

Figure 3: A minimal ROS launch file that deploys a driver node and a higher-level controller node.




Figure 4: Computation Graph of FictiBot as extracted by HAROS. Nodes are shown in white and topics in green.

2.3 The HAROS Framework

HAROS is a framework for static analysis of ROS applications. As of version 3.0, it is capable of extracting Computation Graph models from the source code, although this feature is currently only available for C++ code. This works by the user defining *Configurations* in YAML *project files*, essentially named lists of launch files which are parsed in order. For each launch file, HAROS extracts the set of participating nodes. Afterwards, the source code for each node is parsed and the model extended with the usage of ROS primitives (e.g. subscribing to a topic). As topics are detected, the associated message type is determined and annotated too.

HAROS tries to resolve as many conditions and variables as possible in static time, but, due to the dynamic nature of ROS, this analysis is limited. Nodes may not be launched from a launch file if a condition is not satisfied, and topics are created on demand, as nodes call the ROS API with arguments that may be themselves the result of some computation, or collected from external sources (e.g. parameters). The resulting model marks conditional resources as such (including the respective conditions) when HAROS is unable to fully resolve them. To alleviate this issue, users can provide *hints* when defining a configuration (e.g. topics that a node may subscribe to), which are used by HAROS in the extraction process. Fortunately, conditional topics are not very common in practice [12], and thereby HAROS is able to cover a significant ROS corpus. Figure 4 shows the diagram of our running example, as extracted and depicted by HAROS.

3 PROPERTY-BASED INTEGRATION TESTING FOR ROS

Property-based Testing, as presented in Section 2.1, fits very naturally in unit testing and, with the addition of stateful testing, can even be used to test full components. To leverage PBT at the integration testing level, our approach is to consider the set of integrated

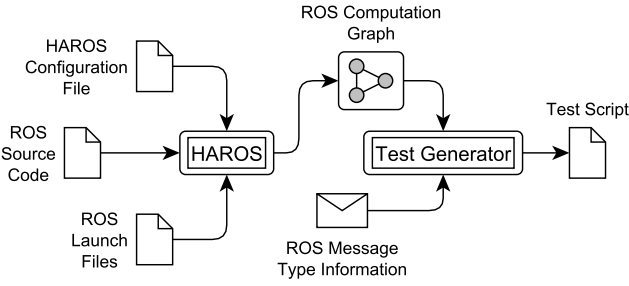


Figure 5: Workflow of the test generation process.

components as a black box, which can then be viewed as if it were a single component with an expected behaviour. In ROS terms, the idea is to choose a set of nodes (e.g. an entire system, or a subset of it), identify subscribed topics to use as inputs, and published topics to take as outputs. The test script should start a ROS node itself, so that communication with the configuration under test can be established. This approach brings forth a major challenge – whereas PBT is often used to test single functions or synchronous systems (the client-server model), we must deal with asynchronous systems (the publisher-subscriber model).

Our main contribution is a test script generator, so that we completely avoid the risk of human error when specifying the architecture under test. We took a Model-based Testing approach to accomplish this goal, using HAROS to construct configuration models from source code. As presented in the previous section, these models include the nodes and topics that make up the Computation Graph, annotated with message type information. The generator’s task is to identify input topics, output topics, and to produce ROS message generators that Hypothesis can use to create data. In order to facilitate the whole process, we extended HAROS to include our test generator as a built-in feature⁹. Figure 5 provides an overview of the workflow we just described.

In the remainder of this section, we provide further details on how the model is used to generate test scripts, as well as how the test scripts work. Each step of the process is illustrated with our FictiBot example.

3.1 Test Script Generation

Before the test script generation process begins, configuration models must be extracted. Thus, the first step, is to specify the intended ROS configurations in the HAROS project files. As mentioned in Section 2.3, this boils down to a list of launch files and optional extraction hints. We assume that the user provides enough specification for the extracted model to be correct. Test settings that play a role in the test generation process are also specified at this stage. Figure 6 shows a HAROS project file using the launch file presented in Figure 3 to construct the model depicted in Figure 4.

Since ROS promotes open and dynamic systems, where participants can join and leave at any time, there is no clear definition of what is a ROS application. Lists of launch files are possibly the best candidate definition, which is why HAROS considers a Configuration as the product of such lists. Providing a list of launch

⁹https://github.com/git-afsantos/haros/blob/dev-make-tests/haros/gen_tests.py

```

1 project: haros_tutorials
2 packages:
3   - fictibot_drivers
4   - fictibot_controller
5 configurations:
6   minimal:
7     launch: [fictibot_controller/launch/minimal.launch]
8     tests: {minimal: {package: fictibot_tests}}
  
```

Figure 6: A minimal HAROS project file for the FictiBot example.

files, instead of a simple list of nodes, is also advantageous for the purposes of test generation, since the generated test script is able to reproduce the configuration as it is in a normal setup. This is an important detail, because the test routine requires launching the same configuration multiple times, as we explain in Section 3.2, and it must be accurate in doing so, including all ROS parameters and evaluated variables.

After the model extraction takes place, the test generation process begins. Identifying the list of participating nodes is straightforward, so the next step is to identify the input and output topics. We treat the configuration under test as a black box. As such, our current approach only considers open topics (topics with subscribers but no publishers, or vice-versa) as candidates for testing. Furthermore, we also discard nodes and topics that the analysis could not fully resolve.

For instance, when testing the whole *minimal* configuration from our example, the only open topic is *teleop_cmd*, with a single subscriber. The corresponding test script would generate a publisher for this topic, and the testing would boil down to trying to make the configuration crash with random messages on this topic. If we tested each node in isolation, however, the generator would create all the corresponding publishers and subscribers for each topic. We considered generating publishers and subscribers for connected topics as well, but this would interfere with the integration we are supposedly testing for correctness.

The final step in the test generation process is to construct ROS message generators for Hypothesis, called *strategies* in Hypothesis terminology. PBT libraries, Hypothesis included, often provide generators for common and basic types of data out of the box, but custom or complex data types must be specified manually. ROS messages follow a well-defined format¹⁰, in which message fields can be of a basic type (numbers and strings), another message type (composition) or lists of one of the previous. As such, the generation of message strategies can be automated by traversing message fields and producing the respective sub-strategy.

Hypothesis already handles basic types and lists, so we just had to provide the corresponding value limits (e.g. a strategy for unsigned 8-bit integer fields that generates values between 0 and 256). To handle composition, we extended the generation algorithm with recursion and caching. Although this is enough to achieve a working implementation, in some cases developers do not expect the full range of values a type may provide, but intend for certain fields to behave like an enumeration (e.g. LED flashing red, green, or turned off). This is a known limitation of ROS, as it supports constants, but

¹⁰<http://wiki.ros.org/msg>


```

1 @strategies.composite
2 def fictibot_msgs_Custom(draw):
3     msg = fictibot_msgs.Custom()
4     msg.state = draw(strategies.sampled_from([0, 1, 2]))
5     msg.pose = draw(geometry_msgs.Pose2D())
6     msg.bumpers = draw(ros_array(ros_bool, length=3))
7     return msg

```

Figure 7: Example of a generated message strategy.

not proper enumerations. Thus, we have implemented support for users to define simple enumerations, using either literal values or constants defined within a message type. This is a way to eliminate many random tests that would provide little value in the end. On the other hand, use of this feature must be pondered, as unexpected inputs can help uncover more faults. Figure 7 shows the automatically generated Hypothesis strategy for a custom message type, illustrating enumerations (the state field), composition (the pose field) and lists (the bumpers field).

3.2 The Test Script

The entry point of a test script performs only a few operations before the test routine starts. First off, in order to use ROS interfaces, the test script has to register itself as a ROS node. Then it creates a TestSettings object and an InternalState object, which it passed down to the test routine. TestSettings objects contain configuration-specific data, such as the required launch files, or the published topics. This is required, since the test routine is generic¹¹. In general, users shall not find the need to edit these, unless they want to ignore a specific topic, for instance.

We cannot determine in advance which properties users might want to verify, so the only way to specify custom properties, besides not crashing, is to edit the test script and add assertions manually. The InternalState class, which is just a template, is meant for this purpose. Users can extend it, and define state variables and callback functions for messages that the test node publishes or receives. The idea behind this is to allow relatively complex properties to be expressed, such as temporal properties (e.g. *a message has arrived within X seconds of a previous one*), or properties that depend on the history of exchanged messages (e.g. *a message field contains monotonic increasing values over time*). Expressing properties this way is also in line with the style advocated in Hypothesis stateful testing. Figure 8 shows the generated InternalState template when testing the ficticontrol node of our example in isolation.

We use Hypothesis' RuleBasedStateMachine, as presented in Section 2.1, to perform stateful testing. Since we are focusing on the publisher-subscriber aspect of ROS, we define two simple operations, *publish_T* and *spin*, where *publish_T* publishes a randomly generated message on topic *T*, and *spin* (using a common terminology in ROS), processes incoming messages and sleeps for a given time, to achieve a certain loop rate. For each input topic, as determined in the test generation process, a *publish* operation is dynamically defined.

Hypothesis will repeatedly try various sequences of *publish_T* and *spin* operations, until a property is falsified or a limit is reached and all properties are deemed true. This implies that, if the first

```

1 class InternalState(object):
2     def __init__(self):
3         self.on_setup()
4     def on_setup(self):
5         pass
6     # event.topic: topic on which a message was sent/received
7     # event.msg: the actual ROS message
8     # event.time: ROS time when the message was sent/received
9     def on_controller_cmd(self, event):
10        pass # callback for the "/controller_cmd" topic
11    def on_stop_cmd(self, event):
12        pass # callback for the "/stop_cmd" topic
13    def on_laser(self, event):
14        pass # callback for the "/laser" topic
15    def on_bumper(self, event):
16        pass # callback for the "/bumper" topic
17    def on_wheel(self, event):
18        pass # callback for the "/wheel" topic

```

Figure 8: Example of a generated template class.

sequence of operations does not run into an error, there must be a way to reset the complete internal state, so the second iteration starts from a deterministic state. This is also one of the big design challenges we faced. It is a simple matter in regular unit tests, where it suffices to define custom *set up* and *tear down* functions that manipulate the state as necessary. When testing a generic set of ROS nodes – which are completely independent processes from the testing script – there is no standard way to reset the internal state of such nodes as required. Thus, we settled on a sensible, although not very performant option, which is to shutdown the whole configuration under test after a successful iteration, so that the configuration can be restarted and likely be in a clean state. It is mostly for this reason that we keep references to the original launch files.

By default, the test node already checks two simple properties. The first property states that all nodes under test must be alive. We make use of internal libraries of ROS to ping other nodes with a timeout. The test fails if any node terminates (due to error or otherwise) or becomes unresponsive. The second property we test for is that the tested ROS interface is stable, i.e. the set of published and subscribed topics should not change during the test run. It is uncommon, in general, to come up with a use case in which closing a topic is a requirement, and so we treat it as an error.

3.3 Preliminary Experiments

To evaluate the usefulness of our framework in a more realistic scenario, we applied it to Kobuki, an education oriented mobile ROS robot that comes out of the box with various different configurations. At the most basic level, it simply runs the *base node*, which is responsible for publishing sensor data and subscribing to velocity commands, similar to the fictibase node from our example. Other configurations, enabled in separate launch files, add various features such as teleoperation, a safety controller, or a random walk controller. We chose the safety controller as our initial test target.

The safety controller subscribes to sensor topics (bumper, wheel drops and cliff detection) and publishes velocity commands. Upon receiving sensor messages, the callback functions simply update the controller's internal state with the new information. Its publishing

¹¹The core of the implementation can be found at <https://github.com/git-afsantos/rosqc>

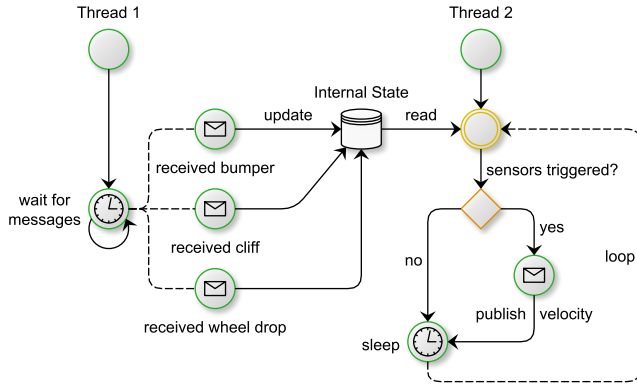


Figure 9: Model of Kobuki's safety controller.

```

1 class InternalState(object):
2     def __init__(self):
3         self.on_setup()
4     def on_setup(self):
5         self.bumper_left = False
6         self.bumper_right = False
7         self.bumper_center = False
8         # other boolean variables ...
9     def on_cmd_vel(self, event):
10        if event.msg.linear.x <= 0:
11            assert (self.bumper_left or self.bumper_right
12                  or self.bumper_center or self.cliff_left
13                  or self.cliff_right or self.cliff_center)
14    def on_events_bumper(self, event):
15        if event.msg.bumper == BumperEvent.LEFT:
16            self.bumper_left =
17                event.msg.state == BumperEvent.PRESSED
18        # repeat for bumper_center and bumper_right
19    def on_events_wheel_drop(self, event):
20        # update internal state
21    def on_events_cliff(self, event):
22        # update internal state

```

Figure 10: Specification of a safety property with internal state.

loop checks the internal state and publishes a zero velocity message if a wheel is dropped, and a negative velocity message (backward movement) when one of the other sensors is triggered. Figure 9 illustrates the node's behaviour.

An immediate property to take out of this setup is that receiving a message of negative velocity implies that the last published message of bumper or cliff sensors should contain an *active* state. For the sake of an illustrative example, we specified that a velocity of zero or less implies a bump or cliff, as seen in Figure 10. Hypothesis has been able to consistently find a minimal counterexample to this property, which is sending a single wheel drop message and waiting for the corresponding zero velocity message. Figure 11 shows a sample of the produced output.

We ran this experiment within a 32-bit virtual machine using a single 2.0 GHz processor and 2 GB of memory. On average, the test runs take about 62 seconds, of which more than 90% of the time is spent setting up and tearing down the ROS configurations (spawning, killing and waiting for processes and network connections),

```

1 FAILED (failures=1)
2 =====
3 Falsifying example
4 -----
5 state = RosRandomTester()
6 state.pub_events__wheel_drop(msg={wheel: 1, state: 1})
7 state.spin()
8 state.teardown()
9 -----
10 Time spent on testing (s): 0.576339435
11 Time spent on sleeping (s): 5.8
12 Time spent setting up (s): 57.640097381

```

Figure 11: Counterexample produced by Hypothesis for an incorrect property.

and only a minimal fraction is spent on the actual testing, publishing and waiting for messages. These results clearly show how detrimental to performance our approach to resetting in-between test iterations is. On the other hand, the fact that Hypothesis is able to find a counterexample shows that our approach is feasible, even if it needs optimisation.

4 RELATED WORK

Automated software testing is very appealing for its promise of finding defects without much effort from the user, but, in [14], Vincenzi et al. show that manual tests tend to be more effective. Even though their work focuses on unit tests for Java programs, we expect the picture to be the same in our context. They state that automated and manual tests have a complementary aspect – whereas manual tests are better to test specific scenarios, automated tests can uncover unexpected faults. We support this view, since, in our case, it would be very hard to test specific scenarios with randomly generated messages.

A proposal for unit testing of publisher-subscriber architectures is presented in [9]. Their proposal is based on Java programs, annotated with preconditions, postconditions and invariants, specified in Linear Temporal Logic. The testing framework mocks the publisher-subscriber infrastructure, so that components can be tested in isolation. Our approach differs in that we are focusing on the ROS infrastructure, and thus we can make use of domain knowledge. Besides, our approach does not require as much (neither as formal) specification in order to function.

When testing stateful systems, most PBT tools generate random transitions based solely on the current state of the model. With the integration of external test case generators, more information can be taken into account, and more meaningful test sequences can be generated. This testing method is presented in [1], and it is an approach that we intend to explore further, since, in many cases, purely random test cases might not make much sense in the context of a ROS application. Our work differs mostly in that we are testing asynchronous distributed systems, our test scripts are automatically generated, and we check architecture-related properties as well.

Despite the importance of testing in robotics, research on the topic is relatively scarce. In [3] the authors present a methodology that bridges the gap between modern software testing techniques and the basic unit testing seen in robotics environments. They apply it to ROS, incorporating simulation environments as substitutes for

real hardware. In [6], this idea is extended to implement Model-based Testing for ROS systems. As is our case, their work is aimed at the integration testing level, but they focus on the navigation and localisation capabilities of mobile robots only. Topological maps, describing the places a robot should be able to move to, are used to generate Timed Automata that simulate the robot's behaviour. Generated tests verify that the robot can move according to the given map, and specific test scenarios can be manually specified, but finding the causes of a test failure is left for the user. While our approach is more user-friendly, in that it produces counterexamples for violated properties, testing of navigation and localisation is possibly one of the hardest features to implement correctly with our approach. As such, our approach might be considered as a complementary test method.

5 CONCLUSIONS AND FUTURE WORK

In this paper we presented our initial approach towards providing a PBT tool for robotic software using ROS. Much of what happens in a ROS system is dependent on messages exchanged between components. This makes it so that specifying component contracts is a given, even if implicitly. PBT brings forth an opportunity to make such contracts explicit, besides providing a mechanism to test them systematically.

Our preliminary results lead us to believe that PBT can be effectively applied to ROS systems, despite the challenges of a dynamic, asynchronous architecture. Testing in the ROS community consists mostly of manual unit tests or test scripts that reproduce a specific scenario in a simulated environment. Introducing automated random testing helps uncover unforeseen faults and provides useful counterexamples for user-specified properties. An interesting point of future work would be to generate specific test cases for each counterexample found.

Another benefit of our approach is to have an automated method of verifying that a ROS configuration produces an expected Computation Graph, where nodes and topics are well integrated. This is a common source of bugs during development, as components are re-used and reconfigured multiple times for different applications. On the other hand, we acknowledge that our proposed tool is no substitute for manual testing, but rather a complementary test method.

Regarding future work directions, there are a few paths in our approach that can be further explored. First off, the languages for user-specified properties and constraints should be extended, for instance to allow other common message constraints, such as ranges for numeric values. There is also the possibility of integrating proper Temporal Logic specifications (e.g. imposing an order on published topics, even though the messages are random). Alternatively, we could integrate external test case generators, using a similar approach to the one presented in [1].

Our current prototype can only generate test scripts for full configurations. If a user wants to test just a subset of the nodes, a new configuration has to be created for that purpose. We intend to work on this matter, either by allowing constraints to be specified, or by generating test scripts for all subsets of a given configuration with either single nodes or nodes connected by at least one topic.

Finally, our current method of resetting configurations between test iterations is lackluster in terms of performance, where a great deal of time is spent in setting up and tearing down processes. This is an immediate target for optimisation, although more performant solutions might not be as generic.

ACKNOWLEDGMENTS

The authors would like to thank the anonymous referees for their valuable comments and helpful suggestions. This work is financed by the ERDF – European Regional Development Fund through the Operational Programme for Competitiveness and Internationalisation - COMPETE 2020 Programme and by National Funds through the Portuguese funding agency, FCT - Fundação para a Ciência e a Tecnologia within project PTDC/CCI-INF/29583/2017 (POCI-01-0145-FEDER-029583).

REFERENCES

- [1] Bernhard K. Aichernig, Silvio Marcovic, and Richard Schumi. 2017. Property-Based Testing with External Test-Case Generators. In *2017 IEEE International Conference on Software Testing, Verification and Validation Workshops (ICSTW)*. 337–346. <https://doi.org/10.1109/ICSTW.2017.62>
- [2] Thomas Arts, John Hughes, Joakim Johansson, and Ulf Wiger. 2006. Testing telecoms software with quviq quickcheck. In *Proceedings of the 2006 ACM SIGPLAN workshop on Erlang*. ACM, 2–10. <https://doi.org/10.1145/1159789.1159792>
- [3] Andreas Bihlmaier and Heinz Wörn. 2014. Robot Unit Testing. In *Simulation, Modeling, and Programming for Autonomous Robots (SIMPAR 2012)*. Springer International Publishing, 255–266. https://doi.org/10.1007/978-3-319-11900-7_22
- [4] Jonathan Bohren, Radu Bogdan Rusu, Edward Gil Jones, Eitan Marder-Eppstein, Caroline Pantofaru, Melonee Wise, Lorenz Mösenlechner, Wim Meeussen, and Stefan Holzer. 2011. Towards autonomous robotic butlers: Lessons learned with the PR2. In *2011 IEEE International Conference on Robotics and Automation*. IEEE, 5568–5575. <https://doi.org/10.1109/ICRA.2011.5980058>
- [5] Koen Claessen and John Hughes. 2000. QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00)*. ACM, 268–279. <https://doi.org/10.1145/351240.351266>
- [6] Juhan Ernits, Evelin Halling, Gert Kanter, and Jüri Vain. 2015. Model-based integration testing of ROS packages: A mobile robot case study. In *2015 European Conference on Mobile Robots (ECMR)*. 1–7. <https://doi.org/10.1109/ECMR.2015.7324210>
- [7] George Fink and Matt Bishop. 1997. Property-based Testing: A New Approach to Testing for Assurance. *SIGSOFT Software Engineering Notes* 22, 4 (July 1997), 74–80. <https://doi.org/10.1145/263244.263267>
- [8] David R. MacIver. 2018. Hypothesis 3.57. <https://github.com/HypothesisWorks/hypothesis>.
- [9] Anton Michlmayr, Pascal Fenkam, and Schahram Dustdar. 2006. Specification-Based Unit Testing of Publish/Subscribe Applications. In *26th IEEE International Conference on Distributed Computing Systems Workshops (ICDCSW'06)*. 34–34. <https://doi.org/10.1109/ICDCSW.2006.103>
- [10] David L. Parnas, A. John Van Schouwen, and Shu Po Kwan. 1990. Evaluation of safety-critical software. *Commun. ACM* 33, 6 (June 1990), 636–648. <https://doi.org/10.1145/78973.78974>
- [11] Morgan Quigley, Ken Conley, Brian P. Gerkey, Josh Faust, Tully Foote, Jeremy Leibs, Rob Wheeler, and Andrew Y. Ng. 2009. ROS: An open-source Robot Operating System. In *ICRA Workshop on Open Source Software*. <https://www.willowgarage.com/sites/default/files/icraoss09-ROS.pdf>
- [12] André Santos, Alcino Cunha, Nuno Macedo, Rafael Arrais, and Filipe Neves dos Santos. 2017. Mining the usage patterns of ROS primitives. In *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 3855–3860. <https://doi.org/10.1109/IROS.2017.8206237>
- [13] André Santos, Alcino Cunha, Nuno Macedo, and Cláudio Lourenço. 2016. A framework for quality assessment of ROS repositories. In *2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 4491–4496. <https://doi.org/10.1109/IROS.2016.7759661>
- [14] Auri M. R. Vincenzi, Tiago Bacheiga, Daniel G. de Oliveira, Simone R. S. de Souza, and José C. Maldonado. 2016. The Complementary Aspect of Automatically and Manually Generated Test Case Sets. In *Proceedings of the 7th International Workshop on Automating Test Case Design, Selection, and Evaluation (A-TEST 2016)*. ACM, 23–30. <https://doi.org/10.1145/2994291.2994295>